

A Block Coordinate Ascent Algorithm for Mean-Variance Optimization

Tengyang Xie^{*1} Bo Liu^{*2} Yangyang Xu³ Mohammad Ghavamzadeh⁴ Yinlam Chow⁵ Daoming Lyu² Daesub Yoon⁶ (* equal contribution)

¹UMass Amherst

²Auburn University

³Rensselaer Polytechnic Institute

⁴Facebook AI Research

⁵Google DeepMind

⁶ETRI

Challenges

Risk-sensitive Reinforcement Learning: How to *manage risk* (e.g., variance of the cumulative reward) for reinforcement learning?

Challenges:

- Computing an unbiased estimate policy gradients with variance related risk criteria usually requires *double sampling* or *multi-time-scale* stochastic approximation algorithm.
- Sample complexity* of existing methods is difficult to analyze.

Our approach: Mean-variance objective function based on its Legendre-Fenchel dual.

Problem Setup

Double sampling issue:

- Mean-variance optimization

$$\max_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}}[R] \quad \text{s.t.} \quad \text{Var}_{\pi_{\theta}}(R) \leq \zeta$$

- Lagrangian relaxation procedure

$$\begin{aligned} J_{\lambda}(\theta) &:= \mathbb{E}_{\pi_{\theta}}[R] - \lambda(\text{Var}_{\pi_{\theta}}(R) - \zeta) \\ &= J(\theta) - \lambda(M(\theta) - J(\theta)^2 - \zeta) \end{aligned}$$

- Computing stochastic gradient ($M(\theta) := \mathbb{E}_{\pi_{\theta}}[R^2]$)

$$\begin{aligned} \nabla_{\theta} J_{\lambda}(\theta_t) &= \nabla_{\theta} J(\theta_t) - \lambda \nabla_{\theta} \text{Var}(R) \\ &= \nabla_{\theta} J(\theta_t) - \lambda(\nabla_{\theta} M(\theta) - 2J(\theta) \nabla_{\theta} J(\theta)) \end{aligned}$$

- We have **no** unbiased estimation of $J(\theta) \nabla_{\theta} J(\theta)$ with a single trajectory.

Multi-time-scale method^[2]:

$$\begin{aligned} \theta_{k+1} &= \theta_k + \alpha_k \left(R^k(\theta_k) - \lambda g'(\tilde{V}_k - b) \left((R^k(\theta_k))^2 - 2R^k(\theta_k) \tilde{J}_k \right) \right) z^k(\theta_k) \\ \tilde{J}_{k+1} &= \tilde{J}_k + \beta_k (R^k(\theta_k) - \tilde{J}_k) \\ \tilde{V}_{k+1} &= \tilde{V}_k + \beta_k \left((R^k(\theta_k))^2 - \tilde{J}_k^2 - \tilde{V}_k \right) \end{aligned}$$

Red terms make it impossible to analyze sample complexity by existing approaches^[1].

Objective: Risk-sensitive reinforcement learning method without the burden of multi-time-scale and double sampling, and with provable sample complexity analysis.

Block Coordinate Reformulation

Reformulation using Legendre-Fenchel dual:

$$\begin{aligned} F_{\lambda}(\theta) &:= \left(J(\theta) + \frac{1}{2\lambda} \right)^2 - M(\theta) \\ &= \max_y \left(2y \left(J(\theta) + \frac{1}{2\lambda} \right) - y^2 \right) - M(\theta) \end{aligned}$$

New optimization problem (standard nonconvex coordinate ascent problem):

$$\max_{\theta, y} \hat{f}_{\lambda}(\theta, y) := 2y \left(J(\theta) + \frac{1}{2\lambda} \right) - y^2 - M(\theta).$$

Mean-Variance Policy Gradient (MVP)

Algorithm 1 Mean-Variance Policy Gradient (MVP)

- Input:** Stepsizes $\{\beta_t^{\theta}\}$ and $\{\beta_t^y\}$, and number of iterations N
Option I: $\{\beta_t^{\theta}\}$ and $\{\beta_t^y\}$ satisfy the Robbins-Monro condition
Option II: β_t^{θ} and β_t^y are set to be constants
- for** episode $t = 1, \dots, N$ **do**
- Generate the initial state $s_1 \sim P_0$
- while** $s_k \neq s^*$ **do**
- Take the action $a_k \sim \pi_{\theta_t}(a|s_k)$ and observe the reward r_k and next state s_{k+1}
- end while**
- Update the parameters

$$\begin{aligned} R_t &= \sum_{k=1}^{\tau_t} r_k \\ \omega_t(\theta_t) &= \sum_{k=1}^{\tau_t} \nabla_{\theta} \ln \pi_{\theta_t}(a_k|s_k) \\ y_{t+1} &= y_t + \beta_t^y \left(2R_t + \frac{1}{\lambda} - 2y_t \right) \\ \theta_{t+1} &= \theta_t + \beta_t^{\theta} \left(2y_{t+1}R_t - (R_t)^2 \right) \omega_t(\theta_t) \end{aligned}$$

- end for**

- Output** \bar{x}_N :

Option I: Set $\bar{x}_N = x_N = [\theta_N, y_N]^{\top}$

Option II: Set $\bar{x}_N = x_z = [\theta_z, y_z]^{\top}$, where z is uniformly drawn from $\{1, 2, \dots, N\}$

Theoretical Analysis of MVP

Asymptotic convergence:

Theorem. Let $\{x_t = (\theta_t, y_t)\}$ be the sequence of the outputs generated by MVP algorithm with Option I. If $\{\beta_t^{\theta}\}$ and $\{\beta_t^y\}$ are time-diminishing real positive sequences satisfying the Robbins-Monro condition, i.e., $\sum_{t=1}^{\infty} \beta_t^{\theta} = \infty$, $\sum_{t=1}^{\infty} (\beta_t^{\theta})^2 < \infty$, $\sum_{t=1}^{\infty} \beta_t^y = \infty$, and $\sum_{t=1}^{\infty} (\beta_t^y)^2 < \infty$, then MVP Algorithm will converge such that

$$\lim_{t \rightarrow \infty} \mathbb{E}[\|\nabla \hat{f}_{\lambda}(x_t)\|_2] = 0.$$

Finite sample analysis:

Theorem. Let the output of the MVP Algorithm be \bar{x}_N following Option II. If $\{\beta_t^{\theta}\}$, $\{\beta_t^y\}$ are constants and satisfies $2\beta_t^{\min} > L(\beta_t^{\max})^2$ for $t = 1, \dots, N$, we have

$$\mathbb{E}[\|\nabla \hat{f}_{\lambda}(\bar{x}_N)\|_2^2] \leq \frac{\hat{f}_{\lambda}^* - \hat{f}_{\lambda}(x_1) + N(\beta_t^{\max})^2 C}{N(\beta_t^{\min} - \frac{L}{2}(\beta_t^{\max})^2)}$$

where $\hat{f}_{\lambda}^* = \max_x \hat{f}_{\lambda}(x)$.

Corollary: The convergence rate of the MVP algorithm with constant stepsizes $\mathcal{O}(1/\sqrt{N})$ implies that the sample complexity $N = \mathcal{O}(1/\varepsilon^2)$ in order to find ε -stationary solution.

Finite-Sample Analysis of Nonconvex Block Stochastic Gradient (BSG) Algorithms

MVP algorithm belongs to the family of nonconvex BSG algorithm

$$\text{objective function: } \min_{x \in \mathbb{R}^n} f(x) = \mathbb{E}_{\xi}[F(x, \xi)]$$

Theorem. Let the output of the nonconvex BSG algorithm be $\bar{x}_N = x_z$. If stepsizes satisfy $2\beta_t^{\min} > L(\beta_t^{\max})^2$ for $t = 1, \dots, N$, then we have

$$\mathbb{E}[\|\nabla f(\bar{x}_N)\|_2^2] \leq \frac{f(x_1) - f^* + \sum_{t=1}^N (\beta_t^{\max})^2 C_t}{\sum_{t=1}^N (\beta_t^{\min} - \frac{L}{2}(\beta_t^{\max})^2)},$$

where $f^* = \min_x f(x)$. $C_t = (1 - \frac{L}{2}\beta_t^{\max}) \sum_{i=1}^b L \sqrt{\sum_{j<i} (G^2 + \sigma^2)} + b \left(AG + \frac{L}{2}\sigma^2 \right)$, where G is the gradient bound, L is the Lipschitz constants, σ is the variance bound.

Corollary: The convergence rate of the nonconvex BSG algorithm with constant stepsizes $\mathcal{O}(1/\sqrt{N})$ implies that the sample complexity $N = \mathcal{O}(1/\varepsilon^2)$ in order to find ε -stationary solution.

Experimental Study




Figure: Empirical results of the distributions of the return (cumulative rewards) random variable. Note that markers only indicate different methods.

	Portfolio Management		American-style Option		Optimal Stopping	
	Mean	Std	Mean	Std	Mean	Std
MVP	29.754	0.325	0.2478	0.00482	-1.4767	0.00456
PG	29.170	1.177	0.2477	0.00922	-1.4769	0.00754
Tamar	28.575	0.857	0.2240	0.00694	-2.8553	0.00415
SGA	29.679	0.658	0.2470	0.00679	-1.4805	0.00583
RCPG	29.340	0.789	0.2447	0.00819	-1.4872	0.00721

Table: Performance Comparison among Algorithms

References

- [1] Gal Dalal, Gagan Thoppe, Balázs Szörényi, and Shie Mannor. Finite sample analysis of two-timescale stochastic approximation with applications to reinforcement learning. In *Proceedings of the 31st Conference On Learning Theory*, pages 1199--1233. PMLR, 06--09 Jul 2018.
- [2] Aviv Tamar, Dotan Di Castro, and Shie Mannor. Policy gradients with variance related risk criteria. In *Proceedings of the twenty-ninth international conference on machine learning*, pages 387--396, 2012.